



BUSINESS UNDERSTANDING

NYC Taxi Demand Prediction



EVALUATION



FUTURE IMPROVEMENTS



OBJECTIVE



DATA COLLECTION



DATA PREPARATION



MODELS



OTHER FEATURES




INTRODUCTION



FEATURE UNDERSTANDING





INTRODUCTION



INTRODUCTION


Iconic history and significance

The concept of taxi medallions

Yellow cabs are concentrated in Manhattan, they serve all five boroughs of NYC.

Market Disruption, Competitive Challenges, Regulatory and Economic Impacts

Operational Efficiency, Economic Viability, Adaptation and Survival



OBJECTIVE



OBJECTIVE

- To analyze historical taxi trip data to understand the demand patterns across different times and locations in New York City.
- To develop a predictive model that can forecast taxi demand, helping in efficient fleet management and reducing passenger wait times.
- To utilize the predictive model to optimize taxi dispatches, thereby increasing operational efficiency and revenue.
- To contribute insights that can help improve urban transportation planning and reduce traffic congestion.
- To integrate both spatial and temporal data to provide precise demand forecasts that account for location-based and time-specific factors.
- To ensure taxis are used efficiently, contributing to a reduction in unnecessary idling and fuel consumption, thereby minimizing the environmental footprint.



BUSINESS UNDERSTANDING



BUSINESS UNDERSTANDING

Industry Shift: The launch of ride-sharing services like Uber in 2011 introduced a new, flexible model of transportation, disrupting the traditional taxi service industry.


Decline in Taxi Use: Challenges faced by the traditional medallion (yellow) cabs with a significant decline in patronage due to the convenience and ease of booking offered by ride-sharing competitors.

Survival Through Adaptation: Emphasis on the need for the traditional taxi industry to adapt by leveraging technology to optimize their operations and compete effectively in the changed landscape.

Data-Driven Dispatching: Introducing the concept of predictive analytics as a solution for smarter dispatching, which could align taxi availability with fluctuating demand patterns.

Strategic Positioning of Cabs: Predictive insights could enable taxi dispatchers to strategically position cabs in high-demand areas, enhancing service responsiveness and customer satisfaction.

Profit Margin Revival: predictive modeling's potential to assist in critical decision-making that could lead to increased operational efficiency, better customer service, and ultimately, the revival of profit margins for the traditional taxi services.



DATA COLLECTION

DATA COLLECTION

The data used for building our solution was collected and provided to the NYC Taxi and Limousine Commission (TLC) by technology providers authorized under the Taxicab & Livery Passenger Enhancement Programs (TPEP/LPEP).

The screenshot shows the NYC Taxi & Limousine Commission website. The header includes the NYC logo and navigation links: About, Passengers, Drivers, Vehicles, Businesses, and TLC Online. Below the header is a search bar and a 'Data and Reports' button. The main content area is titled 'TLC Trip Record Data' and contains two paragraphs of text. The first paragraph describes Yellow and green taxi trip records, and the second paragraph describes For-Hire Vehicle (FHV) trip records. Below the text is an 'ATTENTION!' section with a list of changes to trip record files as of 05/13/2022.

NYC Taxi & Limousine Commission

Русский Translate Text-Size

About Passengers Drivers Vehicles Businesses TLC Online Search

About TLC Data and Reports TLC Initiatives Contact TLC

Data

Pilot Programs

Reports

[TLC Trip Record Data](#)

Request Data

TLC Trip Record Data

Yellow and green taxi trip records include fields capturing pick-up and drop-off dates/times, pick-up and drop-off locations, trip distances, itemized fares, rate types, payment types, and driver-reported passenger counts. The data used in the attached datasets were collected and provided to the NYC Taxi and Limousine Commission (TLC) by technology providers authorized under the Taxicab & Livery Passenger Enhancement Programs (TPEP/LPEP). The trip data was not created by the TLC, and TLC makes no representations as to the accuracy of these data.

For-Hire Vehicle ("FHV") trip records include fields capturing the dispatching base license number and the pick-up date, time, and taxi zone location ID (shape file below). These records are generated from the FHV Trip Record submissions made by bases. Note: The TLC publishes base trip record data as submitted by the bases, and we cannot guarantee or confirm their accuracy or completeness. Therefore, this may not represent the total amount of trips dispatched by all TLC-licensed bases. The TLC performs routine reviews of the records and takes enforcement actions when necessary to ensure, to the extent possible, complete and accurate information.

ATTENTION!

On 05/13/2022, we are making the following changes to trip record files:

1. All files will be stored in the PARQUET format. Please see the 'Working With PARQUET Format' under the Data Dictionaries and MetaData section.
2. Trip data will be published monthly (with two months delay) instead of bi-annually.
3. HVFHV files will now include 17 more columns (please see High Volume FHV Trips Dictionary for details). Additional columns will be added to the old files as well. The earliest date to include additional columns: February 2019.
4. Yellow trip data will now include 1 additional column ('airport_fee', please see Yellow Trips Dictionary for details). The additional column will be added to the old files as well. The earliest

2015	
January <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)	July <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)
February <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)	August <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)
March <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)	September <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)
April <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)	October <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)
May <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)	November <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)
June <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)	December <ul style="list-style-type: none">• Yellow Taxi Trip Records (PARQUET)• Green Taxi Trip Records (PARQUET)• For-Hire Vehicle Trip Records (PARQUET)



FEATURE UNDERSTANDING



FEATURE UNDERSTANDING

temporal span: training - Jan 2015 ; testing - Jan 2016
geographical scope - all five boroughs of NYC

```
In [13]: # Use 2015 data for training and 2016 data for testing
```

```
df_train = df[df["year"] == 2015]
df_test = df[df["year"] == 2016]
```

```
# Create train and test datasets
# Remove year column as it is no longer needed
X_train = df_train.drop(columns=["pickup_count", "year"])
y_train = df_train["pickup_count"]
X_test = df_test.drop(columns=["pickup_count", "year"])
y_test = df_test["pickup_count"]
```

```
print("Train dataset:")
display(X_train)
display(y_train)
```

```
print("Test dataset:")
display(X_test)
display(y_test)
```

Train dataset:

year_sin	day_of_year_cos	time_10min_sin	time_10min_cos	day_of_week_1	day_of_week_2	day_of_week_3	day_of_week_4	day_of_week_5	day_of_week_6
336e-02	0.999852	0.000000	1.000000	False	False	True	False	False	False
336e-02	0.999852	0.043619	0.999048	False	False	True	False	False	False
336e-02	0.999852	0.087156	0.996195	False	False	True	False	False	False
336e-02	0.999852	0.130526	0.991445	False	False	True	False	False	False
336e-02	0.999852	0.173648	0.984808	False	False	True	False	False	False
...
491e-16	1.000000	-0.216440	0.976296	False	False	True	False	False	False
491e-16	1.000000	-0.173648	0.984808	False	False	True	False	False	False
491e-16	1.000000	-0.130526	0.991445	False	False	True	False	False	False
491e-16	1.000000	-0.087156	0.996195	False	False	True	False	False	False
491e-16	1.000000	-0.043619	0.999048	False	False	True	False	False	False


```
In [8]: # Join the weather data with the main dataframe

df = pd.merge(df, df_weather, how="left", on="date")

print("Final dataframe with weather data:")
display(df)
display(df.info())
```

Final dataframe with weather data:

	LocationID	year	day_of_year	time_10min	pickup_count	date	is_holiday	day_of_week	temperature_max	temperature_avg	...	humidity_max
0	1	2015	1	0	0	2015-01-01	True	3	38	33.3	...	46
1	1	2015	1	1	0	2015-01-01	True	3	38	33.3	...	46
2	1	2015	1	2	0	2015-01-01	True	3	38	33.3	...	46
3	1	2015	1	3	0	2015-01-01	True	3	38	33.3	...	46
4	1	2015	1	4	0	2015-01-01	True	3	38	33.3	...	46
...
27856795	265	2016	365	139	4	2016-12-30	False	4	42	38.5	...	82
27856796	265	2016	365	140	1	2016-12-30	False	4	42	38.5	...	82
27856797	265	2016	365	141	4	2016-12-30	False	4	42	38.5	...	82
27856798	265	2016	365	142	1	2016-12-30	False	4	42	38.5	...	82
27856799	265	2016	365	143	5	2016-12-30	False	4	42	38.5	...	82

27856800 rows × 24 columns

VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger	trip_distance	pickup_longitude	pickup_latitude	RateCode	store_and_fwd_flag	dropoff_longitude	dropoff_latitude	payment_type	fare_amount	extra	mta_tax	tip_amount	tolls_amount	improvement_surcharge	total_amount
2	1/15/2015 19:05	1/15/2015 19:23	1	1.59	-73.9939	40.750110	1 N		-73.9748	40.750617	1	12	1	0.5	3.25	0	0.3	17.05
1	1/10/2015 20:33	1/10/2015 20:53	1	3.3	-74.0016	40.72424	1 N		-73.9944	40.759109	1	14.5	0.5	0.5	2	0	0.3	17.8
1	1/10/2015 20:33	1/10/2015 20:43	1	1.8	-73.9633	40.802787	1 N		-73.9518	40.824413	2	9.5	0.5	0.5	0	0	0.3	10.8
1	1/10/2015 20:33	1/10/2015 20:35	1	0.5	-74.0091	40.713817	1 N		-74.0043	40.719985	2	3.5	0.5	0.5	0	0	0.3	4.8
1	1/10/2015 20:33	1/10/2015 20:52	1	3	-73.9712	40.762428	1 N		-74.0042	40.742652	2	15	0.5	0.5	0	0	0.3	16.3
1	1/10/2015 20:33	1/10/2015 20:53	1	9	-73.8744	40.77405	1 N		-73.987	40.758193	1	27	0.5	0.5	6.7	5.33	0.3	40.33
1	1/10/2015 20:33	1/10/2015 20:58	1	2.2	-73.9833	40.726009	1 N		-73.9925	40.74963	2	14	0.5	0.5	0	0	0.3	15.3
1	1/10/2015 20:33	1/10/2015 20:42	3	0.8	-74.0027	40.734142	1 N		-73.995	40.726325	1	7	0.5	0.5	1.66	0	0.3	9.96
1	1/10/2015 20:33	1/10/2015 21:11	3	18.2	-73.783	40.644355	2 N		-73.9876	40.759357	2	52	0	0.5	0	5.33	0.3	58.13
1	1/10/2015 20:33	1/10/2015 20:40	2	0.9	-73.9856	40.767948	1 N		-73.9859	40.759365	1	6.5	0.5	0.5	1.55	0	0.3	9.35
1	1/10/2015 20:33	1/10/2015 20:41	1	0.9	-73.9886	40.723102	1 N		-74.0044	40.728584	1	7	0.5	0.5	1.66	0	0.3	9.96
1	1/10/2015 20:33	1/10/2015 20:43	1	1.1	-73.9938	40.751419	1 N		-73.9674	40.757217	1	7.5	0.5	0.5	1	0	0.3	9.8
1	1/10/2015 20:33	1/10/2015 20:35	1	0.3	-74.0084	40.704376	1 N		-74.0098	40.707725	2	3	0.5	0.5	0	0	0.3	4.3
1	1/10/2015 20:33	1/10/2015 21:03	1	3.1	-73.9739	40.760448	1 N		-73.9973	40.735210	1	19	0.5	0.5	3	0	0.3	23.3
1	1/10/2015 20:33	1/10/2015 20:39	1	1.1	-74.0067	40.731777	1 N		-73.9952	40.739894	2	6	0.5	0.5	0	0	0.3	7.3
2	1/15/2015 19:05	1/15/2015 19:32	1	2.38	-73.9764	40.739810	1 N		-73.984	40.757888	1	16.5	1	0.5	4.38	0	0.3	22.68
2	1/15/2015 19:05	1/15/2015 19:21	5	2.83	-73.9687	40.754245	1 N		-73.9551	40.786857	2	12.5	1	0.5	0	0	0.3	14.3
2	1/15/2015 19:05	1/15/2015 19:28	5	8.33	-73.8631	40.769580	1 N		-73.9527	40.785781	1	26	1	0.5	8.08	5.33	0.3	41.21
2	1/15/2015 19:05	1/15/2015 19:20	1	2.37	-73.9455	40.779422	1 N		-73.9809	40.786083	1	11.5	1	0.5	0	0	0.3	13.3
2	1/15/2015 19:05	1/15/2015 19:20	2	7.13	-73.8745	40.774009	1 N		-73.9524	40.718589	1	21.5	1	0.5	4.5	0	0.3	27.8
2	1/15/2015 19:05	1/15/2015 19:31	1	3.6	-73.9766	40.751895	1 N		-73.9989	40.714595	2	17.5	1	0.5	0	0	0.3	19.3
2	1/15/2015 19:05	1/15/2015 19:10	1	0.89	-73.995	40.745079	1 N		-73.9999	40.734649	1	5.5	1	0.5	1.62	0	0.3	8.92
2	1/15/2015 19:05	1/15/2015 19:10	1	0.96	-74.0009	40.747062	1 N		-74.0036	40.735511	1	5.5	1	0.5	1.3	0	0.3	8.6
2	1/15/2015 19:05	1/15/2015 19:12	2	1.25	-74.0028	40.717891	1 N		-74.0079	40.704219	1	6.5	1	0.5	1.5	0	0.3	9.8
2	1/15/2015 19:05	1/15/2015 19:22	5	2.11	-73.9975	40.736362	1 N		-73.9782	40.761856	1	11.5	1	0.5	2.5	0	0.3	15.8
2	1/15/2015 19:05	1/15/2015 19:14	5	1.15	-73.9592	40.822892	1 N		-73.9522	40.811988	1	7.5	1	0.5	1.7	0	0.3	11.5

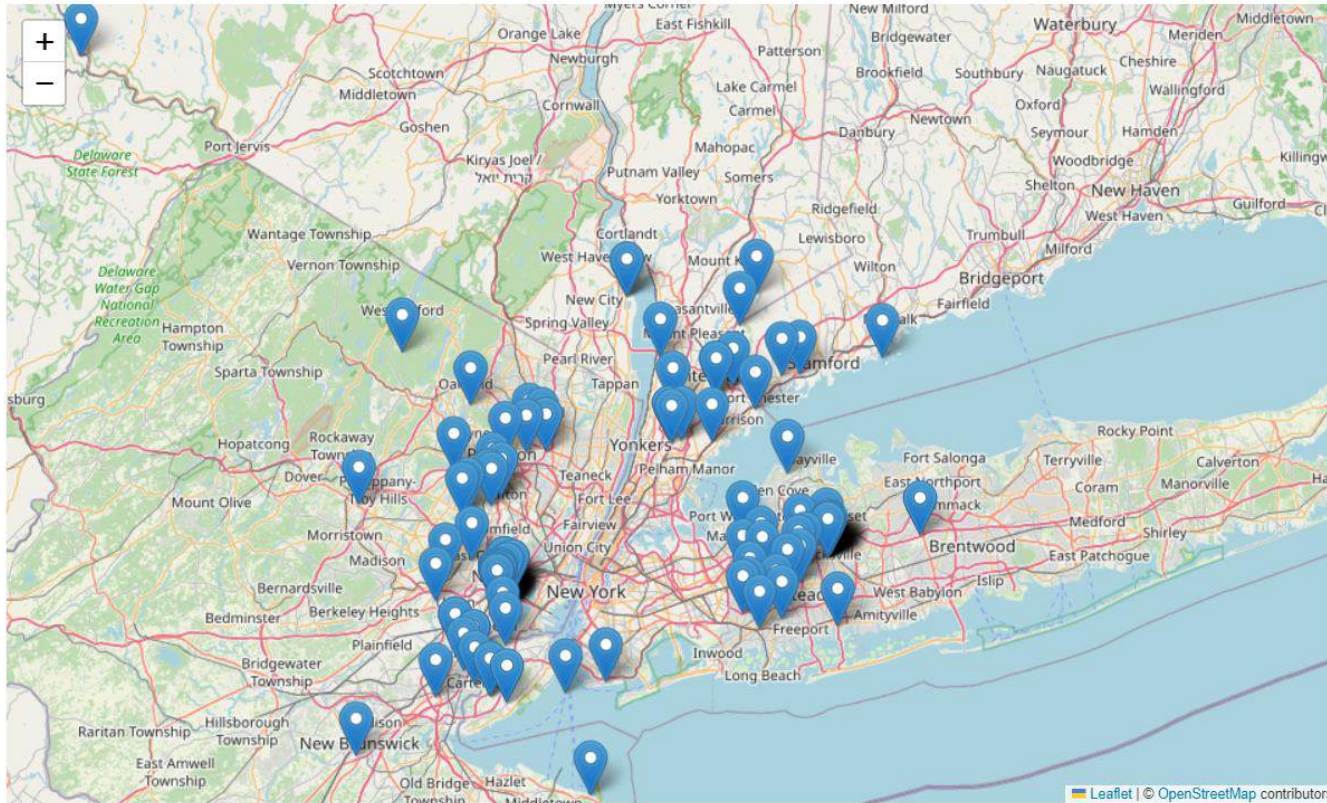


DATA CLEANING

1. Pickup locations:

```
sample_locations = sample_locations.head(1000)
for i,j in sample_locations.iterrows():
    if int(j['pickup_latitude'])!=0:
        folium.Marker(list((j['pickup_latitude'],j['pickup_longitude']))).add_to(map_osm)
map_osm
```

Out[8]:



2. Drop-off locations:

2. Drop-off latitude and longitude.

```
In [9]: #plotting pickup coordinates which are outside the boundary box of newyork
#we will collect all the points outside the bounding the box of newyork city to outlier_location
outlier_locations=month[((month.dropoff_longitude<=-74.15) | (month.dropoff_latitude<=40.5774) | \
                        (month.dropoff_longitude>=-73.7004) | (month.dropoff_latitude>=40.9176))]

map_osm = folium.Map(location=[40.734685, -73.990372], tiles='OpenStreetMap')


sample_locations=outlier_locations.head(10000)
for i,j in sample_locations.iterrows():
    if int(j['dropoff_latitude'])!=0:
        folium.Marker(list((j['dropoff_latitude'],j['dropoff_longitude']))).add_to(map_osm)
map_osm
```

Out[9]:



3.Trip Durations:


According to NYC Taxi & Limousine Commission Regulations the maximum allowed trip duration in a 24 hour interval is 12 hours. The timestamps are converted to unix so as to get duration (trip-time) & speed also pickup-times in unix are used while binning.



```
0 percentile value is -1211.0166666666667
10 percentile value is 3.8333333333333335
20 percentile value is 5.383333333333334
30 percentile value is 6.816666666666666
40 percentile value is 8.3
50 percentile value is 9.95
60 percentile value is 11.866666666666667
70 percentile value is 14.283333333333333
80 percentile value is 17.633333333333333
90 percentile value is 23.45
100 percentile value is 548555.6333333333
```

```
90 percentile value is 23.45
91 percentile value is 24.35
92 percentile value is 25.383333333333333
93 percentile value is 26.55
94 percentile value is 27.933333333333334
95 percentile value is 29.583333333333332
96 percentile value is 31.683333333333334
97 percentile value is 34.466666666666667
98 percentile value is 38.716666666666667
99 percentile value is 46.75
100 percentile value is 548555.6333333333
```


Trip time after removing outliers



4. Speed


```
0 percentile value is 0.0
10 percentile value is 6.409495548961425
20 percentile value is 7.80952380952381
30 percentile value is 8.929133858267717
40 percentile value is 9.98019801980198
50 percentile value is 11.06865671641791
60 percentile value is 12.286689419795222
70 percentile value is 13.796407185628745
80 percentile value is 15.963224893917962
90 percentile value is 20.186915887850468
100 percentile value is 192857142.85714284
```

```
90 percentile value is 20.186915887850468
91 percentile value is 20.91645569620253
92 percentile value is 21.752988047808763
93 percentile value is 22.721893491124263
94 percentile value is 23.844155844155843
95 percentile value is 25.182552504038775
96 percentile value is 26.80851063829787
97 percentile value is 28.84304932735426
98 percentile value is 31.591128254580514
99 percentile value is 35.7513566847558
100 percentile value is 192857142.85714284
```


```
99.0 percentile value is 35.7513566847558
99.1 percentile value is 36.31084727468969
99.2 percentile value is 36.91470054446461
99.3 percentile value is 37.588235294117645
99.4 percentile value is 38.33035714285714
99.5 percentile value is 39.17580340264651
99.6 percentile value is 40.15384615384615
99.7 percentile value is 41.338301043219076
99.8 percentile value is 42.86631016042781
99.9 percentile value is 45.3107822410148
100 percentile value is 192857142.85714284
```

```
In [29]: #avg.speed of cabs in New-York
sum(frame_with_durations_modified["Speed"]) / float(len(frame_with_durations_modified["Speed"]))
```

```
Out[29]: 12.450173996028015
```




Before removal of outliers




After removal of outliers

5. Trip Distance.

0 percentile value is 0.01	90 percentile value is 5.97	Speed at 99.0th percentile: 18.17
10 percentile value is 0.66	91 percentile value is 6.45	Speed at 99.1th percentile: 18.37
20 percentile value is 0.9	92 percentile value is 7.07	Speed at 99.19999999999999th percentile: 18.6
30 percentile value is 1.1	93 percentile value is 7.85	Speed at 99.29999999999998th percentile: 18.83
40 percentile value is 1.39	94 percentile value is 8.72	Speed at 99.39999999999998th percentile: 19.13
50 percentile value is 1.69	95 percentile value is 9.6	Speed at 99.49999999999997th percentile: 19.5
60 percentile value is 2.07	96 percentile value is 10.6	Speed at 99.59999999999997th percentile: 19.96
70 percentile value is 2.6	97 percentile value is 12.1	Speed at 99.69999999999996th percentile: 20.5
80 percentile value is 3.6	98 percentile value is 16.03	Speed at 99.79999999999995th percentile: 21.22
90 percentile value is 5.97	99 percentile value is 18.17	Speed at 99.89999999999995th percentile: 22.57
100 percentile value is 258.9	100 percentile value is 258.9	Speed at 99.99999999999994th percentile: 258.899999




Before removal of outliers



After removal of outliers

6. Total Fare



0 percentile value is -242.55
10 percentile value is 6.3
20 percentile value is 7.8
30 percentile value is 8.8
40 percentile value is 9.8
50 percentile value is 11.16
60 percentile value is 12.8
70 percentile value is 14.8
80 percentile value is 18.3
90 percentile value is 25.8
100 percentile value is 3950611.6

90 percentile value is 25.8
91 percentile value is 27.3
92 percentile value is 29.3
93 percentile value is 31.8
94 percentile value is 34.8
95 percentile value is 38.53
96 percentile value is 42.6
97 percentile value is 48.13
98 percentile value is 58.13
99 percentile value is 66.13
100 percentile value is 3950611.6

Speed at 99.0th percentile: 66.13
Speed at 99.1th percentile: 68.13
Speed at 99.1999999999999th percentile: 69.6
Speed at 99.2999999999998th percentile: 69.6
Speed at 99.3999999999998th percentile: 69.73
Speed at 99.4999999999997th percentile: 69.75
Speed at 99.5999999999997th percentile: 69.76
Speed at 99.6999999999996th percentile: 72.58
Speed at 99.7999999999995th percentile: 75.35
Speed at 99.8999999999995th percentile: 88.27223999984562
Speed at 99.9999999999994th percentile: 3950611.5705955215

Remove all outliers/erronous points.

```
Removing outliers in the month of Jan-2015
```

```
---
```

```
number of pickup records= 12748986
```

```
number of outlier coordinates lying outside NY boundries: 293919
```

```
Number of outliers from the trip time analysis: 23889
```

```
Number of outliers from trip distance analysis: 92597
```

```
Number of outlier from speed analysis: 24473
```

```
Number of outlier from fare analysis: 5275
```

```
Total outliers removed 377910
```

```
---
```

```
fraction of data points that remains after removing outliers 0.9703576425607495
```



DATA PREPARATION




DATA PREPARATION

1. Clustering and segmentation :

Creating cluster using k-means method.

The regions with more number of pickups form smaller and dense clusters whereas regions with lesser number of pickups get connected into larger and loose clusters.

We need to find the optimal number of clusters needed to divide in the geo area, for that we need to check minimum inner cluster should be less and out cluster should be more.



```

On choosing a cluster size of 10
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 2.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 7.0
Min inter-cluster distance = 0.9942822667922672
---
On choosing a cluster size of 20
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 5.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 14.0
Min inter-cluster distance = 0.6444725834028739
---
On choosing a cluster size of 30
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 9.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 21.0
Min inter-cluster distance = 0.47920626820356643
---
On choosing a cluster size of 40
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 11.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 29.0
Min inter-cluster distance = 0.36064577963428435
---
On choosing a cluster size of 50
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 13.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 37.0
Min inter-cluster distance = 0.37726530352041876

```

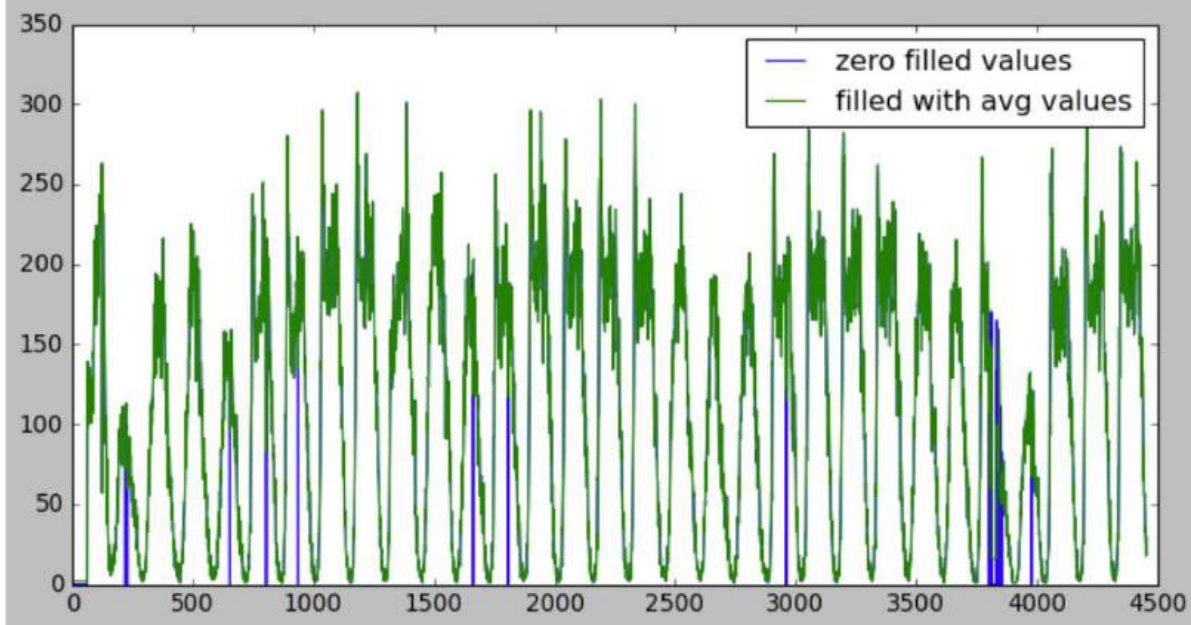


```
Return with trip times..
Remove outliers..
number of pickup records= 10906858
number of outlier coordinates lying outside NY boundries:
Number of outliers from the trip time analysis: 27190
Number of outliers from trip distance analysis: 79742
Number of outlier from speed analysis: 21047
Number of outlier from fare analysis: 4991
Total outliers removed 297784
---
Estimating clusters..
Final groupbying..
```

		trip_distance
pickup_cluster	pickup_bins	
0	63	93
	64	174
	65	208
	66	174



2.SMOOTHING:



```
-----  
for the 3 th cluster number of 10min intavels with zero pickups: 40  
-----  
for the 4 th cluster number of 10min intavels with zero pickups: 270  
-----  
for the 5 th cluster number of 10min intavels with zero pickups: 44  
-----  
for the 6 th cluster number of 10min intavels with zero pickups: 41  
-----  
for the 7 th cluster number of 10min intavels with zero pickups: 35  
-----  
for the 8 th cluster number of 10min intavels with zero pickups: 696  
-----  
for the 9 th cluster number of 10min intavels with zero pickups: 39  
-----  
for the 10 th cluster number of 10min intavels with zero pickups: 37  
-----  
for the 11 th cluster number of 10min intavels with zero pickups: 93  
-----  
for the 12 th cluster number of 10min intavels with zero pickups: 31  
-----  
for the 13 th cluster number of 10min intavels with zero pickups: 28  
-----  
for the 14 th cluster number of 10min intavels with zero pickups: 36  
-----  
for the 15 th cluster number of 10min intavels with zero pickups: 31  
-----  
for the 16 th cluster number of 10min intavels with zero pickups: 56  
-----  
for the 17 th cluster number of 10min intavels with zero pickups: 37  
-----  
for the 18 th cluster number of 10min intavels with zero pickups: 24  
-----  
for the 19 th cluster number of 10min intavels with zero pickups: 40  
-----
```


3. Time Series and Fourier Transformers:

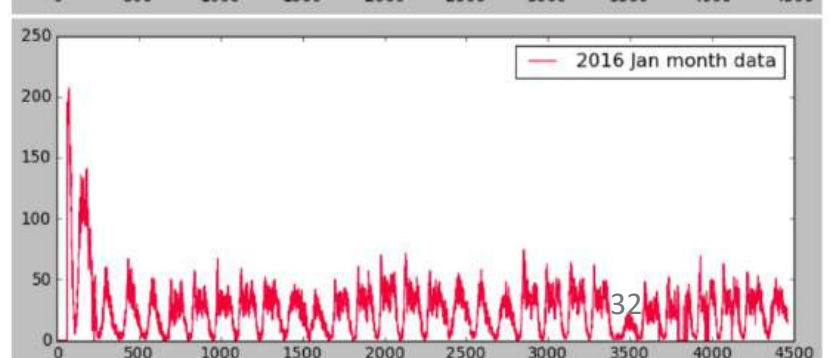

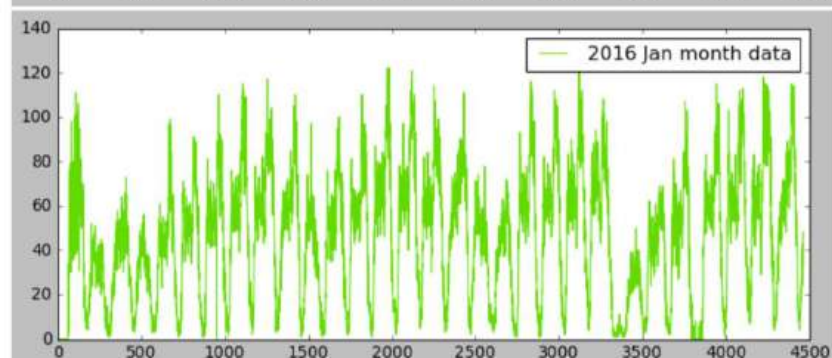






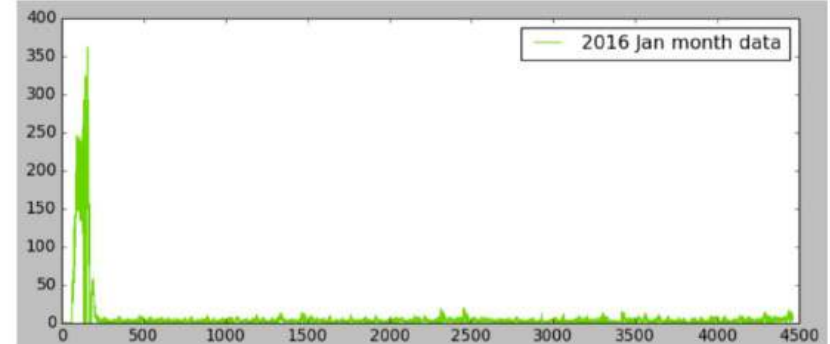


We can see that the number of pickups in a month in every cluster form a repeating pattern.


Fourier transform lets us represent our pattern from time domain (number of pickups per time) to frequency domain (can be viewed as number pickup bins with highest number of pickups).

For each cluster there exists a pattern and using the Fourier transform we can deduce the top frequencies and amplitudes of sine waves which compose our pattern from cluster and use them as features.

The frequencies and amplitudes of a cluster are indicative of demand in that cluster. So they can be fed into the model for prediction of number of pickups.

1. The pattern whose repetition is very high will have a high frequency component and vice versa.
2. There are high frequency in morning and evening time durations as the pick-ups are high during peak hours. The same applies to day and night but with less frequencies.





MODELS



MODELS



BASELINE MODELS



REGRESSOR MODELS




LINEAR REGRESSOR



XGBoost Regressor



RANDOM FOREST



BASELINE MODELS

Baseline model:

1. Using Ratios of the 2016 data to the 2015 data i.e. $\text{Ratio} = P_{2016}/P_{2015}$
2. Using Previous known values of the 2016 data itself to predict the future values

Our 3 choices were **Simple Moving Averages**, **Weighted Moving Averages** and **Exponential Weighted Moving Average**. For each of them, Ratios and Predictions were calculated for P_{2015} and P_{2016} .

Simple Moving Averages:

The First Model used is the Moving Averages Model which uses the previous n values in order to predict the next value

The First Model used is the Moving Averages Model which uses the previous n values in order to predict the next value

Using Ratio Values - $R_t = (R_{t-1} + R_{t-2} + R_{t-3} \dots R_{t-n})/n$

Exponential_moving_average:

Through weighted averaged we have satisfied the analogy of giving higher weights to the latest value and decreasing weights to the subsequent ones but we still do not know which is the correct weighting scheme as there are infinitely many possibilities in which we can assign weights in a non-increasing order and tune the the hyperparameter window-size. To simplify this process we use Exponential Moving Averages which is a more logical way towards assigning weights and at the same time also using an optimal window-size.

In exponential moving averages we use a single hyperparameter alpha which is a value between 0 & 1 and based on the value of the hyperparameter alpha the weights and the window sizes are configured.

$$R'_t = \alpha * R_t + (1 - \alpha) * R'_{t-1}$$


Error Metric Matrix (Forecasting Methods) - MAPE & MSE

Moving Averages (Ratios) -	MAPE: 0.3062982780952648	MSE: 3216.790176971326
Moving Averages (2016 Values) -	MAPE: 0.15087674445032245	MSE: 212.38060035842295

Weighted Moving Averages (Ratios) -	MAPE: 0.31414114299290913	MSE: 2807.2774305555554
Weighted Moving Averages (2016 Values) -	MAPE: 0.14320561490386557	MSE: 192.9197748655914

Exponential Moving Averages (Ratios) -	MAPE: 0.3209817925518331	MSE: 2809.560047043011
Exponential Moving Averages (2016 Values) -	MAPE: 0.1428974374720102	MSE: 191.02174059139784

Please Note:- The above comparisons are made using Jan 2015 and Jan 2016 only




REGRESSOR MODELS

Regressor:

After done with baseline modeling, we have good features which helps us to build a good model prediction

Exponential weighted moving averages gives the best forecasting among the rest. We will use this as a feature while building the regression model along with others we got from data preparation stage.

	ft_5	ft_4	ft_3	ft_2	ft_1	lat	lon	weekday	exp_avg
0	23	26	13	17	21	40.730032	-73.990701	4	19
1	26	13	17	21	39	40.730032	-73.990701	4	33
2	13	17	21	39	27	40.730032	-73.990701	4	28
3	17	21	39	27	42	40.730032	-73.990701	4	37
4	21	39	27	42	53	40.730032	-73.990701	4	48



LINEAR REGRESSOR


Linear regressor:

Linear regression models the relationship between dependent and independent variables. It assumes a linear relationship between the variables. Widely used for prediction and inference in various domain

First Import the LinearRegression class from scikit-learn's linear_model module.
Fit the linear regression model to the training data using the fit method.
Predict the target variable for the testing data using the predict method and round the predicted values.
Predict the target variable for the training data using the predict method and round the predicted values.

```
from sklearn.linear_model import LinearRegression
lr_reg=LinearRegression().fit(df_train, tsne_train_output)

y_pred = lr_reg.predict(df_test)
lr_test_predictions = [round(value) for value in y_pred]
y_pred = lr_reg.predict(df_train)
lr_train_predictions = [round(value) for value in y_pred]
```



RANDOM FOREST

Random Forest:


Random forest is great for maintaining accuracy for a large proportion of data and doesn't allow overfitting if there are too many trees

```
RandomForestRegressor(max_features='sqrt', min_samples_leaf=4,  
                       min_samples_split=3, n_estimators=40, n_jobs=-1)
```

```
y_pred = regr1.predict(df_test)  
rndf_test_predictions = [round(value) for value in y_pred]  
y_pred = regr1.predict(df_train)  
rndf_train_predictions = [round(value) for value in y_pred]
```

```
]: #feature importances based on analysis using random forest  
print (df_train.columns)  
print (regr1.feature_importances_)
```

```
Index(['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1', 'lat', 'lon', 'weekday',  
      'exp_avg'],  
      dtype='object')  
[0.00965372 0.03874404 0.11964895 0.12248967 0.32846239 0.00248174  
 0.00283212 0.00174586 0.37394151]
```

XGBoost Regressor


XGBoost regressor:

Our next choice is XGBoost Regressor which is an implementation of gradient boosted decision trees designed for speed and performance. Training a hyper-parameter tuned Xg-Boost regressor on our train data.

The XGBoost Regressor is part of the XGBoost library, which is specifically designed for regression tasks. XGBoost offers a wide range of hyperparameters that control the model's behavior, including learning rate, number of estimators (trees), maximum tree depth etc.

As in the previous model, we calculated the Mean Absolute Percentage Error


```
{'ft_5': 1049.0, 'ft_4': 667.0, 'ft_3': 834.0, 'ft_2': 885.0,  
'ft_1': 1060.0, 'lat': 715.0, 'lon': 793.0, 'weekday': 232.0,  
'exp_avg': 706.0}
```



EVALUATION

Error Metric Matrix (Tree Based Regression Methods) - MAPE

Baseline Model -	Train: 0.1492063141406475	Test: 0.1541341143087541
Exponential Averages Forecasting -	Train: 0.14112388905282106	Test: 0.1460411665370477
Linear Regression -	Train: 0.14141633700523829	Test: 0.14682461624629067
Random Forest Regression -	Train: 0.09712335885941091	Test: 0.14618996124926825
XgBoost Regression -	Train: 0.13632742704186268	Test: 0.1449294237325879



OTHER FEATURES



OTHER FEATURES - Weather

Weather conditions may impact the demand for taxi services in urban areas like NYC

- Use yearly data to better reflect the seasonal change in the weather
- Total number of ride records from 2015-2016: 277,171,036
- This dataset divides NYC into 263 different zones

	tpep_pickup_datetime	tpep_dropoff_datetime	PULocationID	DOLocationID
0	2015-01-01 00:11:33	2015-01-01 00:16:48	41	166
1	2015-01-01 00:18:24	2015-01-01 00:24:20	166	238
2	2015-01-01 00:26:19	2015-01-01 00:41:06	238	162
3	2015-01-01 00:45:26	2015-01-01 00:53:20	162	263
4	2015-01-01 00:59:21	2015-01-01 01:05:24	236	141
...

Weather dataset

Daily weather of NYC measured by LaGuardia Airport Station

Weather Metrics:

- Temperature
- Dew points
- Humidity
- Wind
- Pressure
- Precipitation

For each metric, maximum, minimum, and average is provided.

Date	Temperature_m	Temperature_av	Temperature_m	Dew_point_māx	Dew_point_avg	Dew_point_min	Humidity_māx
2015-01-01	38	33.3	28	18	10	5	46
2015-01-02	42	38.6	35	22	17.7	15	59
2015-01-03	41	35.5	32	40	29.5	19	100
2015-01-04	56	47.4	41	52	44.2	30	100
2015-01-05	49	35.3	22	28	12.9	3	52

Training and Testing dataset

- For cyclic columns such as time of day, day of year, encode with sine and cosine.
- For day of week, use one-hot encoding.
- 2015 data for training, 2016 data for testing.

LocationID	year	day_of_year	time_10min	pickup_count	date	is_holiday	day_of_week	temperature_max	temperature_avg	...
1	2015	1	0	0	2015-01-01	True	3	38	33.3	...
1	2015	1	1	0	2015-01-01	True	3	38	33.3	...
1	2015	1	2	0	2015-01-01	True	3	38	33.3	...
1	2015	1	3	0	2015-01-01	True	3	38	33.3	...
1	2015	1	4	0	2015-01-01	True	3	38	33.3	...

Ensemble Models

Random Forest:

Mean Squared Error: 47.81

Mean Absolute Percentage Error: 2.25

Gradient Boosting:

Mean Squared Error: 464.62

Mean Absolute Percentage Error: 12.27

Ada Boost:

Mean Squared Error: 1237.62

Mean Absolute Percentage Error: 30.72

Linear Models

Linear Regression:

Mean Squared Error: 625.53

Mean Absolute Percentage Error: 14.96

Ridge:

Mean Squared Error: 625.53

Mean Absolute Percentage Error: 14.96



Bayesian Regression:

Mean Squared Error: 625.47


Mean Absolute Percentage Error: 14.96

Random Forest Regressor


- Estimator: 16
- Mean Squared Error: 47.81
- Mean Absolute Percentage Error: 2.25




Random Forest Regressor




Random Forest Regressor




Linear Regression Models



Linear Regression Models





FUTURE IMPROVEMENTS

Future Improvement

Performance Issue

- Processed dataset takes 4.5GB in RAM (2 years of data)
- For 64GB RAM PC, training crashes if random forest regressor has over 16 estimators

Proposed solutions

- Implement incremental training for the model and feed data in chunks
- Train small models for each location separately

Other features

- Economics

A nighttime photograph of the New York City skyline, viewed from the Manhattan Bridge. The bridge's steel structure and suspension cables are visible in the foreground, leading the eye towards the illuminated skyscrapers of Lower Manhattan. The Hudson River is visible on the left, and the East River on the right. The sky is a deep twilight blue.

Q&A